

# Feed-forward neural networks

CS/DS 541 2020 – Jacob Whitehill (jrwhitehill@wpi.edu)

In Homework 4, you are tasked with implementing backpropagation for a multi-layer neural network. In these notes, a derivation for the special case of 3 layers is given. This may be helpful for implementing the more general case. **Notation:** superscripts denote either the *layer index* or the *example*; subscripts denote an *element* of a matrix or vector.

## 1 Computing partial derivatives in 3-layer neural network


Consider a 3-layer neural network designed to classify images of hand-written digits from the MNIST dataset. Similarly to Homework 4, the input to the network will be a  $28 \times 28$ -pixel image (converted into a 784-dimensional vector); the output will be a vector of 10 probabilities (one for each digit). Specifically, the network you create should implement a function  $f : \mathbb{R}^{784} \rightarrow \mathbb{R}^{10}$ , where:

$$\begin{aligned} \mathbf{z}^{(1)} &= p_1(\mathbf{x}) = \mathbf{W}^{(1)}\mathbf{x} + \mathbf{b}^{(1)} \\ \mathbf{z}^{(2)} &= p_2(\mathbf{h}^{(1)}) = \mathbf{W}^{(2)}\mathbf{h}^{(1)} + \mathbf{b}^{(2)} \\ \mathbf{h}^{(1)} &= f_1(\mathbf{x}) = a_1(p_1(\mathbf{x})) \\ \hat{\mathbf{y}} &= f_2(\mathbf{h}^{(1)}) = a_2(p_2(\mathbf{h}^{(1)})) \\ f(\mathbf{x}) &= f_2(f_1(\mathbf{x})) = a_2(p_2(a_1(p_1(\mathbf{x})))) \end{aligned}$$

For the activation functions  $a_1, a_2$  in your network, use:

$$\begin{aligned} a_1(\mathbf{z}^{(1)}) &= \text{relu}[\mathbf{z}^{(1)}] \\ a_2(\mathbf{z}^{(2)}) &= \text{softmax}(\mathbf{z}^{(2)}) \end{aligned}$$

The network specified above is shown in the figure below:



As usual, the cross-entropy cost function should be

$$J(\mathbf{W}^{(1)}, \mathbf{b}^{(1)}, \mathbf{W}^{(2)}, \mathbf{b}^{(2)}) = -\frac{1}{n} \sum_{i=1}^n \sum_{k=1}^{10} y_k^{(i)} \log \hat{y}_k^{(i)}$$

where  $n$  is the number of examples.

Note that the figure above is rather pedantic in its level of detail compared to most neural network texts, which typically omit  $\mathbf{z}^{(1)}, \mathbf{z}^{(2)}$ . The reason for showing these terms is that it helps to make explicit how to compute the gradient terms for use in gradient descent. For example, to optimize  $\mathbf{W}^{(1)}$ , all we need to compute is the expression below (which is the matrix-product of multiple Jacobian matrices)

$$\frac{\partial J}{\partial \mathbf{W}^{(1)}} = \frac{\partial J}{\partial \hat{\mathbf{y}}} \frac{\partial \hat{\mathbf{y}}}{\partial \mathbf{z}^{(2)}} \frac{\partial \mathbf{z}^{(2)}}{\partial \mathbf{h}^{(1)}} \frac{\partial \mathbf{h}^{(1)}}{\partial \mathbf{z}^{(1)}} \frac{\partial \mathbf{z}^{(1)}}{\partial \mathbf{W}^{(1)}}$$

## 2 Matrix vectorization

As described during lecture, a common technique in matrix calculus is to “vectorize” a matrix into the concatenation of either its rows or its columns. For example, if a matrix is  $m \times n$ , then it can be vectorized into either a  $1 \times (mn)$  row vector or a  $(mn) \times 1$  column vector (depending on what is desired). This is very useful when computing the Jacobian of a vector-valued function with respect to a matrix input. For example,  $\frac{\partial \mathbf{z}^{(1)}}{\partial \mathbf{W}^{(1)}}$  would actually need to be a 3-dimensional tensor since it represents the gradient of a vector with respect to a matrix. However, by vectorizing  $\mathbf{W}^{(1)}$ , we can compute instead a Jacobian with dimensionality  $n \times (km)$  (where  $\mathbf{z}^{(1)}$  has dimensionality  $n$  and  $\mathbf{W}^{(1)}$  is a  $k \times m$  matrix). Using vectorization, we can rewrite the equation above as:

$$\frac{\partial J}{\partial \text{vec}[\mathbf{W}^{(1)}]} = \frac{\partial J}{\partial \hat{\mathbf{y}}} \frac{\partial \hat{\mathbf{y}}}{\partial \mathbf{z}^{(2)}} \frac{\partial \mathbf{z}^{(2)}}{\partial \mathbf{h}^{(1)}} \frac{\partial \mathbf{h}^{(1)}}{\partial \mathbf{z}^{(1)}} \frac{\partial \mathbf{z}^{(1)}}{\partial \text{vec}[\mathbf{W}^{(1)}]}$$

$\begin{matrix} 1 \times 10 & 10 \times 10 & 10 \times 30 & 30 \times 30 & 30 \times (30 * 784) \end{matrix}$

As an example, the equation above shows the dimensions of each Jacobian matrix assuming that the dimensionality of the hidden layer ( $\mathbf{h}^{(1)}$ ) is 30. Once you have computed the gradient of  $J$  with respect to the (vectorized)  $\text{vec}[\mathbf{W}^{(1)}]$ , it is easy to reshape the vectorized gradient  $\frac{dJ}{\partial \text{vec}[\mathbf{W}^{(1)}]}$  ( $1 \times (30 * 784)$ ) into the gradient  $\frac{dJ}{\partial \mathbf{W}^{(1)}}$  ( $30 \times 784$ ).

### 3 Deriving gradient terms for $\mathbf{W}^{(1)}$

Here we derive the gradient update for one particular weight matrix in the 3-layer neural network. The same idea can be used to obtain all the weight matrices and bias vectors.

$$\begin{aligned}
\frac{\partial J}{\partial \hat{\mathbf{y}}} &= \begin{bmatrix} -\frac{y_1}{\hat{y}_1} & \dots & -\frac{y_{10}}{\hat{y}_{10}} \end{bmatrix} \\
\frac{\partial \hat{\mathbf{y}}}{\partial \mathbf{z}^{(2)}} &= \begin{bmatrix} \hat{y}_1(1-\hat{y}_1) & -\hat{y}_1\hat{y}_2 & \dots & -\hat{y}_1\hat{y}_{10} \\ -\hat{y}_2\hat{y}_1 & \hat{y}_2(1-\hat{y}_2) & -\hat{y}_2\hat{y}_3 & \dots & -\hat{y}_2\hat{y}_{10} \\ -\hat{y}_3\hat{y}_1 & -\hat{y}_3\hat{y}_2 & \hat{y}_3(1-\hat{y}_3) & \dots & -\hat{y}_3\hat{y}_{10} \\ \vdots & \dots & & \ddots & \vdots \\ \vdots & \dots & \dots & \dots & \hat{y}_{10}(1-\hat{y}_{10}) \end{bmatrix} \\
\frac{\partial \mathbf{z}^{(2)}}{\partial \mathbf{h}^{(1)}} &= \begin{bmatrix} \frac{\partial(\mathbf{z}^{(2)})_1}{\partial(\mathbf{h}^{(1)})_1} & \frac{\partial(\mathbf{z}^{(2)})_1}{\partial(\mathbf{h}^{(1)})_2} & \dots & \frac{\partial(\mathbf{z}^{(2)})_1}{\partial(\mathbf{h}^{(1)})_{30}} \\ \frac{\partial(\mathbf{z}^{(2)})_2}{\partial(\mathbf{h}^{(1)})_1} & \frac{\partial(\mathbf{z}^{(2)})_2}{\partial(\mathbf{h}^{(1)})_2} & \dots & \frac{\partial(\mathbf{z}^{(2)})_2}{\partial(\mathbf{h}^{(1)})_{30}} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial(\mathbf{z}^{(2)})_{10}}{\partial(\mathbf{h}^{(1)})_1} & \frac{\partial(\mathbf{z}^{(2)})_{10}}{\partial(\mathbf{h}^{(1)})_2} & \dots & \frac{\partial(\mathbf{z}^{(2)})_{10}}{\partial(\mathbf{h}^{(1)})_{30}} \end{bmatrix} \\
&= W_2 \\
\frac{\partial \mathbf{h}^{(1)}}{\partial \mathbf{z}^{(1)}} &= \begin{bmatrix} \text{relu}'[(\mathbf{z}^{(1)})_1] & 0 & \dots & 0 \\ 0 & \text{relu}'[(\mathbf{z}^{(1)})_2] & \dots & 0 \\ 0 & \dots & \ddots & 0 \\ 0 & \dots & \dots & \text{relu}'[(\mathbf{z}^{(1)})_{30}] \end{bmatrix} \\
\text{where } \text{relu}'[x] &= \begin{cases} 1 & \text{if } x > 0 \\ 0 & \text{if } x \leq 0 \end{cases} \\
\frac{\partial \mathbf{z}^{(1)}}{\partial \text{vec}[\mathbf{W}^{(1)}]} &= \begin{bmatrix} \frac{\partial(\mathbf{z}^{(1)})_1}{\partial(\mathbf{W}^{(1)})_{1,1}} & \dots & \frac{\partial(\mathbf{z}^{(1)})_1}{\partial(\mathbf{W}^{(1)})_{1,784}} & \frac{\partial(\mathbf{z}^{(1)})_1}{\partial(\mathbf{W}^{(1)})_{2,1}} & \dots & \frac{\partial(\mathbf{z}^{(1)})_1}{\partial(\mathbf{W}^{(1)})_{2,784}} & \dots & \frac{\partial(\mathbf{z}^{(1)})_1}{\partial(\mathbf{W}^{(1)})_{30,1}} & \dots & \frac{\partial(\mathbf{z}^{(1)})_1}{\partial(\mathbf{W}^{(1)})_{30,784}} \\ \frac{\partial(\mathbf{z}^{(1)})_2}{\partial(\mathbf{W}^{(1)})_{1,1}} & \dots & \frac{\partial(\mathbf{z}^{(1)})_2}{\partial(\mathbf{W}^{(1)})_{1,784}} & \frac{\partial(\mathbf{z}^{(1)})_2}{\partial(\mathbf{W}^{(1)})_{2,1}} & \dots & \frac{\partial(\mathbf{z}^{(1)})_2}{\partial(\mathbf{W}^{(1)})_{2,784}} & \dots & \frac{\partial(\mathbf{z}^{(1)})_2}{\partial(\mathbf{W}^{(1)})_{30,1}} & \dots & \frac{\partial(\mathbf{z}^{(1)})_2}{\partial(\mathbf{W}^{(1)})_{30,784}} \\ \vdots & \vdots & \vdots & \vdots & \ddots & \vdots & \vdots & \vdots & \vdots & \vdots \\ \frac{\partial(\mathbf{z}^{(1)})_{30}}{\partial(\mathbf{W}^{(1)})_{1,1}} & \dots & \frac{\partial(\mathbf{z}^{(1)})_{30}}{\partial(\mathbf{W}^{(1)})_{1,784}} & \frac{\partial(\mathbf{z}^{(1)})_{30}}{\partial(\mathbf{W}^{(1)})_{2,1}} & \dots & \frac{\partial(\mathbf{z}^{(1)})_{30}}{\partial(\mathbf{W}^{(1)})_{2,784}} & \dots & \frac{\partial(\mathbf{z}^{(1)})_{30}}{\partial(\mathbf{W}^{(1)})_{30,1}} & \dots & \frac{\partial(\mathbf{z}^{(1)})_{30}}{\partial(\mathbf{W}^{(1)})_{30,784}} \end{bmatrix} \\
&= \begin{bmatrix} \mathbf{x}^\top & \mathbf{0} & \mathbf{0} & \dots & \mathbf{0} \\ \mathbf{0} & \mathbf{x}^\top & \mathbf{0} & \dots & \mathbf{0} \\ \mathbf{0} & \mathbf{0} & \mathbf{x}^\top & \dots & \mathbf{0} \\ \mathbf{0} & \mathbf{0} & \mathbf{0} & \ddots & \mathbf{0} \\ \mathbf{0} & \mathbf{0} & \mathbf{0} & \mathbf{0} & \mathbf{x}^\top \end{bmatrix}
\end{aligned}$$

## 4 Analytical simplification

Now, watch what happens when we multiply the matrices together (and compare with Algorithm 6.4 from the book – note that below we are using no regularization and hence  $\lambda = 0$ ):

$$\begin{aligned}
\frac{\partial J}{\partial \hat{\mathbf{y}}} \frac{\partial \hat{\mathbf{y}}}{\partial \mathbf{z}^{(2)}} &= \begin{bmatrix} \hat{y}_1 - y_1 & \cdots & \hat{y}_{10} - y_{10} \end{bmatrix} \\
&= (\hat{\mathbf{y}}^\top - \mathbf{y}^\top) \\
\frac{\partial J}{\partial \hat{\mathbf{y}}} \frac{\partial \hat{\mathbf{y}}}{\partial \mathbf{z}^{(2)}} \frac{\partial \mathbf{z}^{(2)}}{\partial \mathbf{h}^{(1)}} &= (\hat{\mathbf{y}}^\top - \mathbf{y}^\top) \mathbf{W}^{(2)} \\
\frac{\partial J}{\partial \hat{\mathbf{y}}} \frac{\partial \hat{\mathbf{y}}}{\partial \mathbf{z}^{(2)}} \frac{\partial \mathbf{z}^{(2)}}{\partial \mathbf{h}^{(1)}} \frac{\partial \mathbf{h}^{(1)}}{\partial \mathbf{z}^{(1)}} &= (\hat{\mathbf{y}}^\top - \mathbf{y}^\top) \mathbf{W}^{(2)} \begin{bmatrix} \text{relu}'[(\mathbf{z}^{(1)})_1] & 0 & \cdots & 0 \\ 0 & \text{relu}'[(\mathbf{z}^{(1)})_2] & \cdots & 0 \\ 0 & \cdots & \ddots & 0 \\ 0 & \cdots & \cdots & \text{relu}'[(\mathbf{z}^{(1)})_{30}] \end{bmatrix} \\
&= (\hat{\mathbf{y}}^\top - \mathbf{y}^\top) \mathbf{W}^{(2)} \odot \begin{bmatrix} \text{relu}'[(\mathbf{z}^{(1)})_1] & \cdots & \text{relu}'[(\mathbf{z}^{(1)})_{30}] \end{bmatrix} \\
&= (\hat{\mathbf{y}}^\top - \mathbf{y}^\top) \mathbf{W}^{(2)} \odot \text{relu}' \left[ \mathbf{z}^{(1)\top} \right] \\
&\quad \text{where } \odot \text{ represents Hadamard product.}
\end{aligned}$$

To obtain the final expression, first define

$$\begin{aligned}
\mathbf{g}^\top &= (\hat{\mathbf{y}}^\top - \mathbf{y}^\top) \mathbf{W}^{(2)} \odot \text{relu}' \left[ \mathbf{z}^{(1)\top} \right] \\
&= \begin{bmatrix} g_1 & g_2 & \cdots & g_{30} \end{bmatrix}
\end{aligned}$$

Then:

$$\begin{aligned}
\frac{\partial J}{\partial \hat{\mathbf{y}}} \frac{\partial \hat{\mathbf{y}}}{\partial \mathbf{z}^{(2)}} \frac{\partial \mathbf{z}^{(2)}}{\partial \mathbf{h}^{(1)}} \frac{\partial \mathbf{h}^{(1)}}{\partial \mathbf{z}^{(1)}} \frac{\partial \mathbf{z}^{(1)}}{\partial \text{vec}[\mathbf{W}^{(1)}]} &= \mathbf{g}^\top \begin{bmatrix} \mathbf{x}^\top & \mathbf{0} & \mathbf{0} & \cdots & \mathbf{0} \\ \mathbf{0} & \mathbf{x}^\top & \mathbf{0} & \cdots & \mathbf{0} \\ \mathbf{0} & \mathbf{0} & \mathbf{x}^\top & \cdots & \mathbf{0} \\ \mathbf{0} & \mathbf{0} & \mathbf{0} & \ddots & \mathbf{0} \\ \mathbf{0} & \mathbf{0} & \mathbf{0} & \cdots & \mathbf{x}^\top \end{bmatrix} \\
&= \begin{bmatrix} g_1 \mathbf{x}^\top & g_2 \mathbf{x}^\top & \cdots & g_{30} \mathbf{x}^\top \end{bmatrix}
\end{aligned}$$

We can now “unvectorize”  $\frac{\partial J}{\partial \text{vec}[\mathbf{W}^{(1)}]}$  into a matrix with 30 rows and 784 columns:

$$\begin{aligned}
\frac{\partial J}{\partial \mathbf{W}^{(1)}} &= \begin{bmatrix} g_1 \mathbf{x}^\top \\ g_2 \mathbf{x}^\top \\ \vdots \\ g_{30} \mathbf{x}^\top \end{bmatrix} \\
&= \mathbf{g} \mathbf{x}^\top \quad (\text{outer product})
\end{aligned}$$

## 5 Deriving gradient terms for $\mathbf{b}^{(1)}$

Let’s derive  $\frac{\partial J}{\partial \mathbf{b}^{(1)}}$ . Since  $\mathbf{b}^{(1)}$  is just a vector rather than a matrix, we won’t have to vectorize it.

$$\frac{\partial J}{\partial \mathbf{b}^{(1)}} = \frac{\partial J}{\partial \hat{\mathbf{y}}} \frac{\partial \hat{\mathbf{y}}}{\partial \mathbf{z}^{(2)}} \frac{\partial \mathbf{z}^{(2)}}{\partial \mathbf{h}^{(1)}} \frac{\partial \mathbf{h}^{(1)}}{\partial \mathbf{z}^{(1)}} \frac{\partial \mathbf{z}^{(1)}}{\partial \mathbf{b}^{(1)}}$$

$\begin{matrix} 1 \times 10 & 10 \times 10 & 10 \times 30 & 30 \times 30 & 30 \times 30 \end{matrix}$

Most of the terms above are the same as for  $\frac{\partial J}{\partial \mathbf{W}^{(1)}}$  except the last one:

$$\begin{aligned}\frac{\partial \mathbf{z}^{(1)}}{\partial \mathbf{b}^{(1)}} &= \begin{bmatrix} 1 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & \ddots & 0 \\ 0 & 0 & 0 & \dots & 1 \end{bmatrix} \\ &= \mathbf{I}\end{aligned}$$

which is diagonal since each  $(\mathbf{z}^{(1)})^{(i)}$  depends only on the  $i$ th component of  $\mathbf{b}^{(1)}$ . We can simplify the whole sequence of matrix products to become:

$$\frac{\partial J}{\partial \mathbf{b}^{(1)}} = \mathbf{g}$$

## 6 Deriving gradient terms for $\mathbf{W}^{(2)}$

It can be shown, using the same principle as illustrated above, that

$$\frac{\partial J}{\partial \mathbf{W}^{(2)}} = (\hat{\mathbf{y}} - \mathbf{y})\mathbf{h}^{(1)\top}$$

## 7 Forward and backward propagation

As shown above, when computing the gradients of the cost function  $J$  with respect to each weight matrix and bias term, there is a lot of redundant computation (e.g.,  $(\hat{\mathbf{y}} - \mathbf{y})$ ,  $\mathbf{g}$ ) – even after simplifying each gradient analytically. Rather than re-compute these matrices/vectors repeatedly, we can compute *all* gradient terms using a two-pass sweep: (1) forward propagation and (2) backward propagation.

### 7.1 Forward propagation

The job of forward propagation (through layers  $1, 2, \dots, l$ ) is to determine the value of each hidden layer  $\mathbf{h}^{(1)}, \mathbf{h}^{(2)}, \dots$  (in the 3-layer neural network there is only one hidden layer, but this is only a special case) and also the final output layer  $\hat{\mathbf{y}}$ :

1. Let  $\mathbf{h}^{(0)} = \mathbf{x}$ .
2. For  $i = 1, \dots, l - 1$  (where  $l$  is the number of layers, including input and output layers):
  - (a) Let  $\mathbf{z}^{(i)} = p_i(\mathbf{h}^{(i-1)}) = \mathbf{W}^{(i)}\mathbf{h}^{(i-1)} + \mathbf{b}^{(i)}$ .
  - (b) Let  $\mathbf{h}^{(i)} = a_i(\mathbf{z}^{(i)})$ .
3. Let  $\hat{\mathbf{y}} = \mathbf{h}^{(l-1)}$ .

### 7.2 Back(ward)-prop(agation)

Once  $\hat{\mathbf{y}}$  has been computed, we can compute  $(\hat{\mathbf{y}} - \mathbf{y})$ , which is known as the *error* since it expresses the difference between the ground-truth and the network's current predictions. By iteratively multiplying – layer-by-layer ( $l - 1, l - 2, \dots, 2, 1$ ) – by the derivative of each layer's activation function as well as its weight matrices, we can compute and update our vector  $\mathbf{g}$ :

1. Let  $\mathbf{g}^\top = \frac{\partial J}{\partial \hat{\mathbf{y}}}$ .
2. Let  $\mathbf{h}^{(l-1)} = \hat{\mathbf{y}}$ .

3. For  $i = l - 1, \dots, 1$ :

- (a) Let  $\mathbf{g}^\top \leftarrow \mathbf{g}^\top \frac{\partial \mathbf{h}^{(i)}}{\partial \mathbf{z}^{(i)}}$ .
- (b) Let  $\frac{\partial J}{\partial \mathbf{b}^{(i)}} = \mathbf{g}$ .
- (c) Let  $\frac{\partial J}{\partial \mathbf{W}^{(i)}} = \mathbf{g} \mathbf{h}^{(i-1)\top}$ .
- (d) Let  $\mathbf{g}^\top \leftarrow \mathbf{g}^\top \frac{\partial \mathbf{z}^{(i)}}{\partial \mathbf{h}^{(i-1)}} = \mathbf{g}^\top \mathbf{W}^{(i)}$ .

### 7.2.1 Special case for NNs with softmax outputs with cross-entropy loss

In the special case such as Homework 4 in which (1) the network is trained using cross-entropy loss and (2) the network has a softmax output layer, the algorithm above can be simplified slightly. In particular, we can skip the first step (where we initialize  $\mathbf{g}$  to  $\frac{\partial J}{\partial \hat{\mathbf{y}}}$ ); instead, in the *first* iteration of the loop we replace step (a) with simply “Let  $\mathbf{g}^\top = \hat{\mathbf{y}} - \mathbf{y}$ ”. Where does this result come from? You proved this in Homework 3, problem 2.

## 8 Convexity and parameter initialization

In contrast to logistic or softmax regression, the cost function  $J$  of a non-linear 3-layer neural network is generally **not convex**. This means that, depending on the initial values for the weight matrices and bias terms, stochastic gradient descent (SGD) could end up in a different place. While weight initialization is an important topic in its own right, for now let’s just examine a few important cases:

1. Suppose that you initialized all the parameters ( $\mathbf{W}^{(1)}, \mathbf{W}^{(2)}, \mathbf{b}^{(1)}, \mathbf{b}^{(2)}$ ) in the 3-layer neural network in the example above to be 0. How will SGD perform?
  - If  $\mathbf{W}^{(1)}$  and  $\mathbf{b}^{(1)}$  are 0, then  $\mathbf{h}^{(1)} = 0$ , and hence  $\mathbf{z}^{(2)} = 0$ , so that  $\hat{\mathbf{y}}$  is just the uniform distribution over  $c$  classes. Moreover, if  $\text{relu}'$  is defined so that  $\text{relu}'[0] = 0$ , then the gradients with respect to all the weights and bias vectors will be 0, and the network will never learn.
2. Suppose that you initialized all the parameters ( $\mathbf{W}^{(2)}, \mathbf{b}^{(2)}$ ) to be 0 (but  $\mathbf{W}^{(1)}, \mathbf{b}^{(1)}$  are non-zero). How will SGD perform?
  - In this case,  $\hat{\mathbf{y}}$  will *initially* just be the uniform distribution. However, assuming  $\mathbf{h}^{(1)}$  is initially non-zero (and non-uniform),  $\mathbf{W}^{(2)}$  and  $\mathbf{b}^{(2)}$  can still improve. Very soon, the outputs of the network will be non-uniform, and the network can train to good accuracy.