

# CS/DS 541: Class 3

Jacob Whitehill

# Hyperparameter tuning

# Hyperparameter tuning

- The values we **optimize** when training a machine learning model — e.g.,  $\mathbf{w}$  and  $b$  for linear regression — are the **parameters** of the model.
- There are also values related to the training process itself — e.g., learning rate  $\epsilon$ , batch size  $\tilde{n}$ , regularization strength  $\alpha$  — which are the **hyperparameters** of training.

# Hyperparameter tuning

- Both the parameters and hyperparameters can have a huge impact on model performance on test data.
- When estimating the performance of a trained model, it is important to tune both kinds of parameters in a principled way:
  - Training/validation/testing sets
  - Double cross-validation

# Training/validation/testing sets

- In an application domain with a large dataset (e.g., 100K examples), it is common to partition it into three subsets:
  - Training (typically 70-80%): optimization of parameters
  - Validation (typically 5-10%): tuning of hyperparameters
  - Testing (typically 5-10%): evaluation of the final model
- For comparison with other researchers' methods, this partition should be fixed.

# Training/validation/testing sets

- Hyperparameter tuning works as follows:
  1. For each hyperparameter configuration  $h$ :
    - Train the parameters on the **training** set using  $h$ .
    - Evaluate the model on the **validation** set.
    - If performance is better than what we got with the best  $h$  so far ( $h^*$ ), then save  $h$  as  $h^*$ .
  2. Evaluate the model trained with  $h^*$  on the **testing** set.  
(You can train either on training data, or on training+validation data).

# Training/validation/testing sets

- Question: Which machine did we evaluate?

# Training/validation/testing sets

- Question: Which machine did we evaluate?
- Answer: The one trained in step 2:

2. Evaluate the model trained with  $h^*$  on the **testing** set.  
(You can train either on training data, or on training+validation data).



# Cross-validation

- When working with smaller datasets, cross-validation is commonly used so that we can use **all** data for training.
- Assume we already know the best hyperparameters  $h^*$ .
- We partition the data into  $k$  folds of equal sizes.
- Over  $k$  iterations, we train on  $(k-1)$  folds and test on the remaining fold.
- We then compute the average accuracy over the  $k$  testing folds.

# Cross-validation

- #  $D$ =dataset,  $k$ =# folds,  $h$ =hyperparameter configuration.

CrossValidation ( $D, k, h$ ):

Partition  $D$  into  $k$  folds  $F_1, \dots, F_k$

For  $i = 1, \dots, k$ :

$test = F_i$

$train = D \setminus F_i$

Train the model on  $train$  using  $h$

$acc[i] = \text{Evaluate NN on test}$

$A = \text{Avg}[acc]$

return  $A$

# Training/validation/testing sets

- Question: To what machine does the reported accuracy correspond?

# Training/validation/testing sets

- Question: To what machine does the reported accuracy correspond?
- Answer: None of them! Cross-validation gives the *expected* accuracy of a classifier that is trained on  $(k-1)/k$  of the data.

# Training/validation/testing sets

- Question: To what machine does the reported accuracy correspond?
- Answer: None of them! Cross-validation gives the *expected* accuracy of a classifier that is trained on  $(k-1)/k$  of the data.
- However, we can train another model  $M$  using  $h^*$  on the entire dataset, and then report  $A$  as its accuracy.
- Since  $M$  is trained on more data than any of the cross-validation models, its *expected* accuracy should be  $\geq A$ .

# Cross-validation

- But how do we find the best hyperparameters  $h^*$  for each fold?
- The typical approach is to use **double cross-validation**, i.e.:
  - For each of the  $k$  “outer” folds, run cross-validation in an “inner” loop to determine the best hyperparameter configuration  $h^*$  for the  $k$ th fold.

# Double cross-validation

- #  $D$ =dataset,  $k$ =# folds,  $H$ =set of hyperparameter configurations.

DoubleCrossValidation ( $D, k, H$ ):

Partition  $D$  into  $k$  folds  $F_1, \dots, F_k$

**For each of the  $k$  “outer” folds,  
run cross-validation in an “inner”  
loop to determine the best  
hyperparameter configuration  $h^*$   
for the  $k$ th fold.**

**At the end of the procedure,  
which “machine” are you  
evaluating?**

return ...

## For your reference...

CrossValidation ( $D, k, h$ ):

Partition  $D$  into  $k$  folds  $F_1, \dots, F_k$

For  $i = 1, \dots, k$ :

$test = F_i$

$train = D \setminus F_i$

Train the model on  $train$  using  $h$

$acc[i] = \text{Evaluate NN on test}$

$A = \text{Avg}[acc]$

return  $A$

# Double cross-validation

- #  $D$ =dataset,  $k$ =# folds,  $H$ =set of hyperparameter configurations.

DoubleCrossValidation ( $D, k, H$ ):

Partition  $D$  into  $k$  folds  $F_1, \dots, F_k$

For  $i = 1, \dots, k$ :

$test = F_i$

$train = D \setminus F_i$

$A^* = -\infty$

For  $h$  in  $H$ :

$A = \text{CrossValidation}(train, k, h)$

if  $A > A^*$ :

$A^* = A$

$h^* = h$

Train the model on  $train$  using  $h^*$

$accs[i] = \text{Evaluate the model on } test$

$A = \text{Avg}[accs]$

return  $A$

## For your reference...

CrossValidation ( $D, k, h$ ):

Partition  $D$  into  $k$  folds  $F_1, \dots, F_k$

For  $i = 1, \dots, k$ :

$test = F_i$

$train = D \setminus F_i$

Train the model on  $train$  using  $h$

$acc[i] = \text{Evaluate NN on test}$

$A = \text{Avg}[acc]$

return  $A$



# Training/validation/testing sets

- Question: To what machine does the reported accuracy correspond?

# Training/validation/testing sets

- Question: To what machine does the reported accuracy correspond?
- Answer: None of them!
- In contrast to (single) cross-validation, it's not obvious how to train a model  $M$  with accuracy  $\geq A$ .
- One strategy: return an **ensemble** model whose output is the average of the  $k$  models' predictions...but this is rarely done.

# Subject independence

- In many machine learning settings, the data are not completely independent from each other — they are *linked* in some way.
- Example:
  - Predict multiple grades for each student based on their Canvas clickstream features (# logins, # forum posts, etc.).

# Subject independence

- We could partition the data into folds in different ways:
  - We could randomize across all the data.
  - However, if grades are correlated within each student, then one (or more) training folds can leak information about the testing fold.

	Quiz 1	Quiz 2	Quiz 3
Student 1	45	48	42
Student 2	96	93	93
Student 3	86	86	87
Student 4	10	30	50

# Subject independence

- We could partition the data into folds in different ways:
- Alternatively, we can **stratify** across students, i.e., no student appears in more than 1 fold.
- With this partition, the cross-validation accuracy estimates the model's performance on a subject *not* used for training.

	Quiz 1	Quiz 2	Quiz 3
Student 1	45	48	42
Student 2	96	93	93
Student 3	86	86	87
Student 4	10	30	50

# Optimization of ML models

# Optimization of ML models

- Gradient descent is guaranteed to converge to a local minimum (eventually) if the learning rate is small enough relative to the steepness of  $f$ .
- A function  $f : \mathbb{R}^m \rightarrow \mathbb{R}$  is Lipschitz-continuous if:
$$\exists L : \forall \mathbf{x}, \mathbf{y} \in \mathbb{R}^m : \|f(\mathbf{x}) - f(\mathbf{y})\|_2 \leq L \|\mathbf{x} - \mathbf{y}\|_2$$
- $L$  is essentially an upper bound on the absolute slope of  $f$ .


# Optimization of ML models

- Gradient descent is guaranteed to converge to a local minimum (eventually) if the learning rate is small enough relative to the steepness of  $f$ .
- A function  $f : \mathbb{R}^m \rightarrow \mathbb{R}$  is Lipschitz-continuous if:
$$\exists L : \forall \mathbf{x}, \mathbf{y} \in \mathbb{R}^m : \|f(\mathbf{x}) - f(\mathbf{y})\|_2 \leq L \|\mathbf{x} - \mathbf{y}\|_2$$
- $L$  is essentially an upper bound on the absolute slope of  $f$ .
- For learning rate  $\epsilon \leq \frac{1}{L}$ , gradient descent will converge to a local minimum linearly, i.e., the error is  $O(1/k)$  in the iterations  $k$ .



# Optimization of ML models

- With linear regression, the cost function  $f_{\text{MSE}}$  has a single local minimum w.r.t. the weights  $\mathbf{w}$ :




- As long as our learning rate is small enough, we will find **the optimal  $\mathbf{w}$** .

# Optimization: what can go wrong?

- In general ML and DL models, optimization is usually not so simple, due to:


# Optimization: what can go wrong?

- In general ML and DL models, optimization is usually not so simple, due to:
  1. Presence of multiple local minima




# Optimization: what can go wrong?

- In general ML and DL models, optimization is usually not so simple, due to:
  2. Bad initialization of the weights  $\mathbf{w}$ .




# Optimization: what can go wrong?

- In general ML and DL models, optimization is usually not so simple, due to:
  1. Learning rate is too large.
  2. Learning rate is too large.
  3. Learning rate is too small.




# Optimization: what can go wrong?

- In general ML and DL models, optimization is usually not so simple, due to:
  1. Learning rate is too large.
  2. Learning rate is too big.
  3. Learning rate is too small.



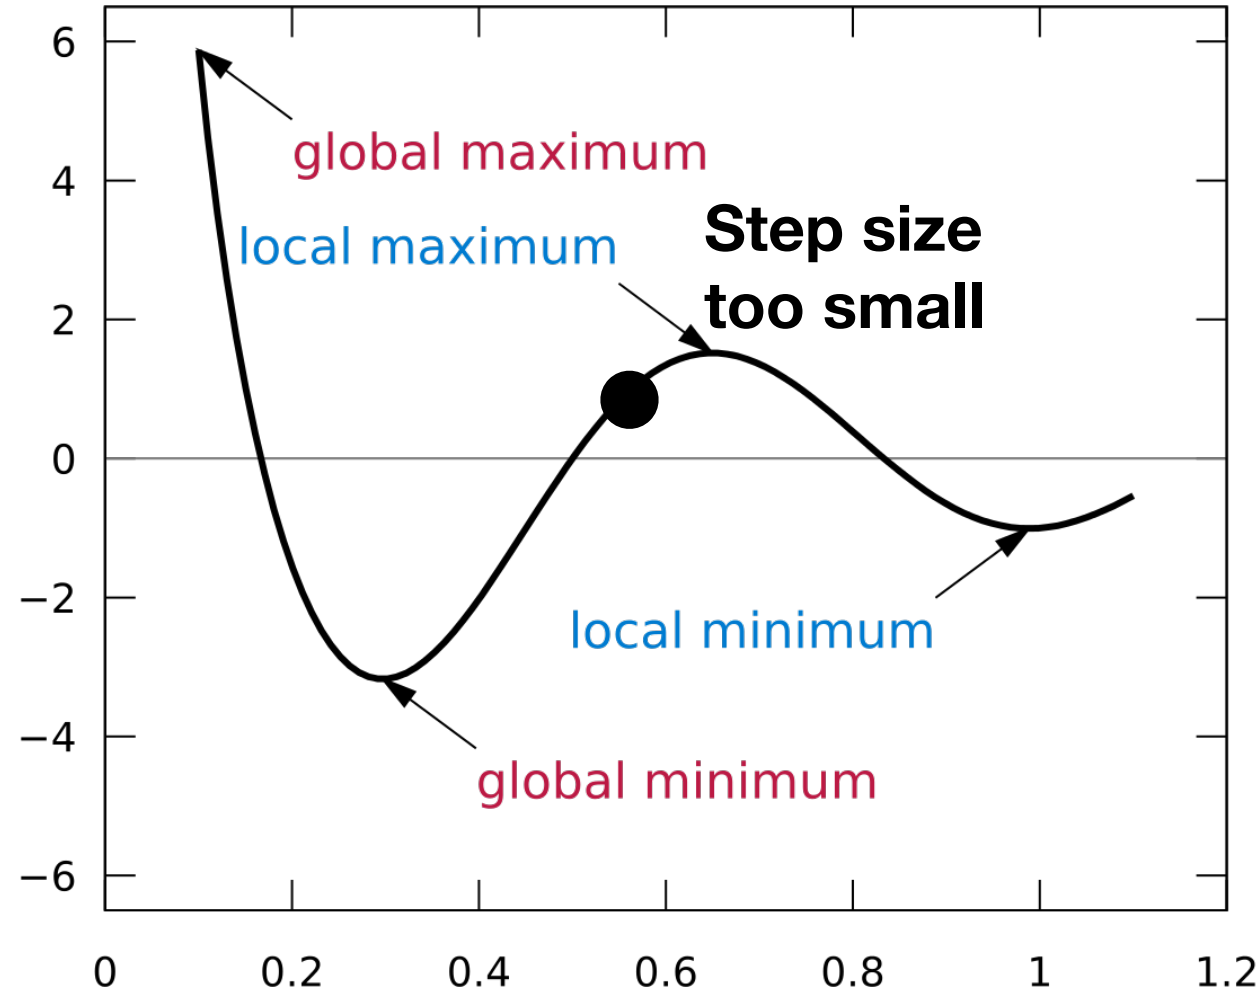
# Optimization: what can go wrong?

- In general ML and DL models, optimization is usually not so simple, due to:
  1. Learning rate is too large.
  2. Learning rate is too large.
  3. Learning rate is too small.



# Optimization: what can go wrong?


- In general ML and DL models, optimization is usually not so simple, due to:
  1. Learning rate is too large.
  2. Learning rate is too large.
  3. Learning rate is too small.






# Optimization: what can go wrong?

- In general ML and DL models, optimization is usually not so simple, due to:
  4. Learning rate is too big.



# Optimization: what can go wrong?

- In general ML and DL models, optimization is usually not so simple, due to:
  4. Learning rate is too big.




# Optimization: what can go wrong?

- In general ML and DL models, optimization is usually not so simple, due to:


4. Learning rate is too big.

●  
(off the chart)




# Optimization: what can go wrong?

- With multidimensional weight vectors, badly chosen learning rates can cause more subtle problems.
- Consider the cost  $f$  whose level sets are shown below:




# Optimization: what can go wrong?

- With multidimensional weight vectors, badly chosen learning rates can cause more subtle problems.
- Gradient descent guides the search along the direction of steepest decrease in  $f$ .




# Optimization: what can go wrong?

- With multidimensional weight vectors, badly chosen learning rates can cause more subtle problems.
- Gradient descent guides the search along the direction of steepest decrease in  $f$ .



# Optimization: what can go wrong?

- With multidimensional weight vectors, badly chosen learning rates can cause more subtle problems.
- Gradient descent guides the search along the direction of steepest decrease in  $f$ .



# Optimization: what can go wrong?

- With multidimensional weight vectors, badly chosen learning rates can cause more subtle problems.
- But what if the level sets are ellipsoids instead of spheres?






# Optimization: what can go wrong?

- With multidimensional weight vectors, badly chosen learning rates can cause more subtle problems.
- But what if the level sets are ellipsoids instead of spheres?
- If we are lucky, we still converge quickly.




# Optimization: what can go wrong?

- With multidimensional weight vectors, badly chosen learning rates can cause more subtle problems.
- But what if the level sets are ellipsoids instead of spheres?
- If we are lucky, we still converge quickly.



# Optimization: what can go wrong?

- With multidimensional weight vectors, badly chosen learning rates can cause more subtle problems.
- But what if the level sets are ellipsoids instead of spheres?
- If we are lucky, we still converge quickly.



# Optimization: what can go wrong?

- With multidimensional weight vectors, badly chosen learning rates can cause more subtle problems.
- But what if the level sets are ellipsoids instead of spheres?
- If we are unlucky, convergence is very slow.




# Optimization: what can go wrong?

- With multidimensional weight vectors, badly chosen learning rates can cause more subtle problems.
- But what if the level sets are ellipsoids instead of spheres?
- If we are unlucky, convergence is very slow.




# Optimization: what can go wrong?

- With multidimensional weight vectors, badly chosen learning rates can cause more subtle problems.
- But what if the level sets are ellipsoids instead of spheres?
- If we are unlucky, convergence is very slow.




# Optimization: what can go wrong?

- With multidimensional weight vectors, badly chosen learning rates can cause more subtle problems.
- But what if the level sets are ellipsoids instead of spheres?
- If we are unlucky, convergence is very slow.



# Optimization: what can go wrong?

- With multidimensional weight vectors, badly chosen learning rates can cause more subtle problems.
- But what if the level sets are ellipsoids instead of spheres?
- If we are unlucky, convergence is very slow.







# Convexity

# Convex ML models

- Linear regression has a loss function that is **convex**.
- With a convex function  $f$ , *every local minimum is also a global minimum*.



convex




non-convex

- Convex functions are ideal for conducting gradient descent.

# Convexity in 1-d


- How can we tell if a 1-d function  $f$  is convex?



- What property of  $f$  ensures there is only one local minimum?

# Convexity in 1-d


- How can we tell if a 1-d function  $f$  is convex?



- What property of  $f$  ensures there is only one local minimum?
- From left to right, the slope of  $f$  *never decreases*.

# Convexity in 1-d


- How can we tell if a 1-d function  $f$  is convex?



- What property of  $f$  ensures there is only one local minimum?
- From left to right, the slope of  $f$  *never decreases*.  
==> the derivative of the slope is always non-negative.

# Convexity in 1-d

- How can we tell if a 1-d function  $f$  is convex?



- What property of  $f$  ensures there is only one local minimum?
- From left to right, the slope of  $f$  *never decreases*.  
==> the derivative of the slope is always non-negative.  
==> the second derivative of  $f$  is always non-negative.

# Convexity in higher dimensions

- For higher-dimensional  $f$ , convexity is determined by the the Hessian of  $f$ .

$$\mathbf{H}[f] = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1 \partial x_1} & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_m} \\ \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_m \partial x_1} & \cdots & \frac{\partial^2 f}{\partial x_m \partial x_m} \end{bmatrix}$$

- For  $f : \mathbb{R}^m \rightarrow \mathbb{R}$ ,  $f$  is convex if the Hessian matrix is positive semi-definite for *every* input  $\mathbf{x}$ .

# Positive semi-definite

- Positive semi-definite is the matrix analog of being “non-negative”.
- A real symmetric matrix **A** is **positive semi-definite (PSD)** if (equivalent conditions):



# Positive semi-definite

- Positive semi-definite is the matrix analog of being “non-negative”.
- A real symmetric matrix **A** is **positive semi-definite (PSD)** if (equivalent conditions):
  - All its eigenvalues are  $\geq 0$ .
  - If **A** happens to be diagonal, then its eigenvalues are the diagonal elements.

# Positive semi-definite

- Positive semi-definite is the matrix analog of being “non-negative”.
- A real symmetric matrix **A** is **positive semi-definite (PSD)** if (equivalent conditions):
  - All its eigenvalues are  $\geq 0$ .
    - If **A** happens to be diagonal, then its eigenvalues are the diagonal elements.
  - For every vector **v**:  $\mathbf{v}^T \mathbf{A} \mathbf{v} \geq 0$ 
    - Therefore: If there exists *any* vector **v** such that  $\mathbf{v}^T \mathbf{A} \mathbf{v} < 0$ , then **A** is *not* PSD.

# Example

- Suppose  $f(x, y) = 3x^2 + 2y^2 - 2$ .

- Then the first derivatives are:  $\frac{\partial f}{\partial x} = 6x$      $\frac{\partial f}{\partial y} = 4y$


- The Hessian matrix is therefore:

$$\mathbf{H} = \begin{bmatrix} \frac{\partial^2 f}{\partial x \partial x} & \frac{\partial^2 f}{\partial x \partial y} \\ \frac{\partial^2 f}{\partial y \partial x} & \frac{\partial^2 f}{\partial y \partial y} \end{bmatrix} = \begin{bmatrix} 6 & 0 \\ 0 & 4 \end{bmatrix}$$

- Notice that  $\mathbf{H}$  for this  $f$  does not depend on  $(x, y)$ .
- Also,  $\mathbf{H}$  is a diagonal matrix (with 6 and 4 on the diagonal). Hence, the eigenvalues are just 6 and 4. Since they are both non-negative, then  $f$  is convex.

# Example

- Graph of  $f(x, y) = 3x^2 + 2y^2 - 2$ :



# Exercise


- Recall: if  $\mathbf{H}$  is the Hessian of  $f$ , then  $f$  is convex if — at every  $(x,y)$ , we can show (equivalently):
  - $\mathbf{v}^T \mathbf{H} \mathbf{v} \geq 0$  for every  $\mathbf{v}$
  - All eigenvalues of  $\mathbf{H}$  are non-negative.
- Which of the following function(s) are convex?
  - $x^2 + y + 5$
  - $x^4 + xy + x^2$

# Exercise

- Recall: if  $\mathbf{H}$  is the Hessian of  $f$ , then  $f$  is convex if — at every  $(x,y)$ , we can show (equivalently):
  - $\mathbf{v}^T \mathbf{H} \mathbf{v} \geq 0$  for every  $\mathbf{v}$
  - All eigenvalues of  $\mathbf{H}$  are non-negative.
- Which of the following function(s) are convex?
  - $x^2 + y + 5$      $\mathbf{H} = \begin{bmatrix} 2 & 0 \\ 0 & 0 \end{bmatrix}$
  - $x^4 + xy + x^2$

# Exercise

- Recall: if  $\mathbf{H}$  is the Hessian of  $f$ , then  $f$  is convex if — at *every*  $(x,y)$ , we can show (equivalently):
  - $\mathbf{v}^T \mathbf{H} \mathbf{v} \geq 0$  for *every*  $\mathbf{v}$
  - All eigenvalues of  $\mathbf{H}$  are non-negative.
- Which of the following function(s) are convex?
  - $x^2 + y + 5$      $\mathbf{H} = \begin{bmatrix} 2 & 0 \\ 0 & 0 \end{bmatrix}$     Eigenvalues are 2, 0  $\Rightarrow$  PSD.
  - $x^4 + xy + x^2$



# Exercise

- Recall: if  $\mathbf{H}$  is the Hessian of  $f$ , then  $f$  is convex if — at every  $(x,y)$ , we can show (equivalently):
  - $\mathbf{v}^T \mathbf{H} \mathbf{v} \geq 0$  for every  $\mathbf{v}$
  - All eigenvalues of  $\mathbf{H}$  are non-negative.
- Which of the following function(s) are convex?
  - $x^2 + y + 5$
  - $x^4 + xy + x^2$      $\mathbf{H} = \begin{bmatrix} 12x^2 + 2 & 1 \\ 1 & 0 \end{bmatrix}$



# Exercise

- Recall: if  $\mathbf{H}$  is the Hessian of  $f$ , then  $f$  is convex if — at every  $(x,y)$ , we can show (equivalently):

- $\mathbf{v}^T \mathbf{H} \mathbf{v} \geq 0$  for every  $\mathbf{v}$
- All eigenvalues of  $\mathbf{H}$  are non-negative.

- Which of the following function(s) are convex?

- $x^2 + y + 5$

$$x = 1$$

- $x^4 + xy + x^2$      $\mathbf{H} = \begin{bmatrix} 12x^2 + 2 & 1 \\ 1 & 0 \end{bmatrix}$      $\mathbf{v} = \begin{bmatrix} -1 \\ 15 \end{bmatrix}$      $\mathbf{v}^T \mathbf{H} \mathbf{v} = -16$   
Not PSD.

# Exercise

- Recall: if  $\mathbf{H}$  is the Hessian of  $f$ , then  $f$  is convex if — at every  $(x,y)$ , we can show (equivalently):

- $\mathbf{v}^T \mathbf{H} \mathbf{v} \geq 0$  for every  $\mathbf{v}$
- All eigenvalues of  $\mathbf{H}$  are non-negative.

- Which of the following function(s) are convex?


- $x^2 + y + 5$

$$x = 1$$

- $x^4 + xy + x^2$       $\mathbf{H} = \begin{bmatrix} 12x^2 + 2 & 1 \\ 1 & 0 \end{bmatrix}$       $\mathbf{v} = \begin{bmatrix} -1 \\ 15 \end{bmatrix}$

$$\mathbf{v}^T \mathbf{H} \mathbf{v} = -16$$

Not PSD.



# Convexity of linear regression

- How do we know linear regression is a convex ML model?
- First, recall that, for any matrices **A**, **B** that can be multiplied:
  - $(\mathbf{AB})^T = \mathbf{B}^T \mathbf{A}^T$

# Convexity of linear regression

- How do we know linear regression is a convex ML model?
- Next, recall the gradient and Hessian of  $f_{\text{MSE}}$  (for linear regression):

$$\begin{aligned}\nabla_{\mathbf{w}} f_{\text{MSE}} &= \mathbf{X}(\hat{\mathbf{y}} - \mathbf{y}) \\ &= \mathbf{X}(\mathbf{X}^{\top} \mathbf{w} - \mathbf{y}) \\ \mathbf{H} &= \mathbf{X}\mathbf{X}^{\top}\end{aligned}$$

# Convexity of linear regression

- How do we know linear regression is a convex ML model?
- Next, recall the gradient and Hessian of  $f_{\text{MSE}}$  (for linear regression):

$$\begin{aligned}\nabla_{\mathbf{w}} f_{\text{MSE}} &= \mathbf{X}(\hat{\mathbf{y}} - \mathbf{y}) \\ &= \mathbf{X}(\mathbf{X}^{\top} \mathbf{w} - \mathbf{y}) \\ \mathbf{H} &= \mathbf{X}\mathbf{X}^{\top}\end{aligned}$$

- For any vector  $\mathbf{v}$ , we have:

$$\begin{aligned}\mathbf{v}^{\top} \mathbf{X}\mathbf{X}^{\top} \mathbf{v} &= (\mathbf{X}^{\top} \mathbf{v})^{\top} (\mathbf{X}^{\top} \mathbf{v}) \\ &\geq 0\end{aligned}$$

# Convex ML models

- Prominent convex models in ML include linear regression, logistic regression, softmax regression, and support vector machines (SVM).
- However, models in deep learning are generally not convex.
- Much DL research is devoted to how to optimize the weights to deliver good generalization performance.

# Non-spherical loss functions

# Non-spherical loss functions


- As described previously, loss functions that are non-spherical can make hill climbing via gradient descent more difficult:






# Curvature

- The problem is that gradient descent only considers slope (1st-order effect), i.e., how  $f$  changes with  $\mathbf{w}$ .




# Curvature

- The problem is that gradient descent only considers slope (1st-order effect), i.e., how  $f$  changes with  $\mathbf{w}$ .
- The gradient does not consider how the slope *itself* changes with  $\mathbf{w}$  (2nd-order effect).



# Curvature

- The problem is that gradient descent only considers slope (1st-order effect), i.e., how  $f$  changes with  $\mathbf{w}$ .
- The gradient does not consider how the slope *itself* changes with  $\mathbf{w}$  (2nd-order effect).
- The higher-order effects determine the curvature of  $f$ .



# Curvature

- For linear regression with cost  $f_{\text{MSE}}$ ,

$$f_{\text{MSE}}(\mathbf{w}) = \frac{1}{2n} (\mathbf{X}^T \mathbf{w} - \mathbf{y})^T (\mathbf{X}^T \mathbf{w} - \mathbf{y})$$

the Hessian is:

$$\mathbf{H}[f](\mathbf{w}) = \frac{1}{n} \mathbf{X} \mathbf{X}^T$$

- Hence,  $\mathbf{H}$  is constant and is proportional to the (uncentered) **auto-covariance matrix** of  $\mathbf{X}$ .

$$\mathbb{E}[(\mathbf{X} - \mathbb{E}[\mathbf{X}])(\mathbf{X} - \mathbb{E}[\mathbf{X}])^T]$$

# Curvature

- For linear regression with cost  $f_{\text{MSE}}$ ,

$$f_{\text{MSE}}(\mathbf{w}) = \frac{1}{2n} (\mathbf{X}^T \mathbf{w} - \mathbf{y})^T (\mathbf{X}^T \mathbf{w} - \mathbf{y})$$

the Hessian is:

$$\mathbf{H}[f](\mathbf{w}) = \frac{1}{n} \mathbf{X} \mathbf{X}^T$$

- In other words, the curvature depends solely on the matrix of training data  $\mathbf{X}$ .

$$\mathbb{E}[(\mathbf{X} - \mathbb{E}[\mathbf{X}])(\mathbf{X} - \mathbb{E}[\mathbf{X}])^T]$$

# Curvature

- To accelerate optimization of the weights, we can either:
  - Alter the cost function by transforming the input data.
  - Change our optimization method to account for the curvature.

# Feature transformations

# Whitening transformations

- Gradient descent works best when the level sets of the cost function are spherical.
- We can “spherize” the input features using a **whitening transformation**, which makes the auto-covariance matrix equal the identity matrix  $\mathbf{I}$ .
- We compute this transformation on the training data, and then apply it to both training and testing data.



# Whitening transformations

- We can find a whitening transform  $\mathbf{T}$  as follows:
  - Let the auto-covariance<sup>\*</sup> of our training data be  $\mathbf{XX}^T$ .

<sup>\*</sup> (uncentered).

# Whitening transformations

- We can find a whitening transform  $\mathbf{T}$  as follows:
  - Let the auto-covariance\* of our training data be  $\mathbf{XX}^\top$ .
  - We can write its eigendecomposition as:

$$\mathbf{XX}^\top \Phi = \Phi \Lambda$$

where  $\Phi$  is the matrix of eigenvectors and  $\Lambda$  is the corresponding diagonal matrix of eigenvalues.

\* (uncentered).

# Whitening transformations

- We can find a whitening transform  $\mathbf{T}$  as follows:
  - Let the auto-covariance\* of our training data be  $\mathbf{XX}^\top$ .
  - We can write its eigendecomposition as:

$$\mathbf{XX}^\top \Phi = \Phi \Lambda$$

where  $\Phi$  is the matrix of eigenvectors and  $\Lambda$  is the corresponding diagonal matrix of eigenvalues.

- For real-valued features,  $\mathbf{XX}^\top$  is real and symmetric; hence,  $\Phi$  is orthonormal. Also,  $\Lambda$  is non-negative.

\* (uncentered).

# Whitening transformations

- We can find a whitening transform  $\mathbf{T}$  as follows:
- Therefore, we can multiply both sides by  $\Phi^\top$ :

$$\mathbf{X}\mathbf{X}^\top \Phi = \Phi \Lambda$$

$$\Phi^\top \mathbf{X}\mathbf{X}^\top \Phi = \Phi^\top \Phi \Lambda = \Lambda$$

# Whitening transformations

- We can find a whitening transform  $\mathbf{T}$  as follows:

- Therefore, we can multiply both sides by  $\Phi^\top$ :

$$\mathbf{X}\mathbf{X}^\top \Phi = \Phi \Lambda$$

$$\Phi^\top \mathbf{X}\mathbf{X}^\top \Phi = \Phi^\top \Phi \Lambda = \Lambda$$

- Since  $\Lambda$  is diagonal and non-negative, we can easily compute  $\Lambda^{-\frac{1}{2}}$ .

# Whitening transformations

- We can find a whitening transform  $\mathbf{T}$  as follows:

- Therefore, we can multiply both sides by  $\Phi^\top$ :

$$\mathbf{X}\mathbf{X}^\top \Phi = \Phi \Lambda$$

$$\Phi^\top \mathbf{X}\mathbf{X}^\top \Phi = \Phi^\top \Phi \Lambda = \Lambda$$

- Since  $\Lambda$  is diagonal and non-negative, we can easily compute  $\Lambda^{-\frac{1}{2}}$ .

- We then multiply both sides (2x) to obtain  $\mathbf{I}$  on the RHS.

$$\Lambda^{-\frac{1}{2}^\top} \Phi^\top \mathbf{X}\mathbf{X}^\top \Phi \Lambda^{-\frac{1}{2}} = \Lambda^{-\frac{1}{2}^\top} \Lambda \Lambda^{-\frac{1}{2}}$$

# Whitening transformations

- We can find a whitening transform  $\mathbf{T}$  as follows:

- Therefore, we can multiply both sides by  $\Phi^\top$ :

$$\mathbf{X}\mathbf{X}^\top \Phi = \Phi \Lambda$$

$$\Phi^\top \mathbf{X}\mathbf{X}^\top \Phi = \Phi^\top \Phi \Lambda = \Lambda$$

- Since  $\Lambda$  is diagonal and non-negative, we can easily compute  $\Lambda^{-\frac{1}{2}}$ .

- We then multiply both sides (2x) to obtain  $\mathbf{I}$  on the RHS.

$$\Lambda^{-\frac{1}{2}\top} \Phi^\top \mathbf{X}\mathbf{X}^\top \Phi \Lambda^{-\frac{1}{2}} = \Lambda^{-\frac{1}{2}\top} \Lambda \Lambda^{-\frac{1}{2}}$$

$$\left( \Lambda^{-\frac{1}{2}\top} \Phi^\top \mathbf{X} \right)^\top \left( \Lambda^{-\frac{1}{2}\top} \Phi^\top \mathbf{X} \right) = \mathbf{I}$$

# Whitening transformations

- We have thus derived a transform  $\mathbf{T} = \Lambda^{-\frac{1}{2}} \Phi^\top$  such that the (uncentered) auto-covariance of the transformed data  $\tilde{\mathbf{X}} = \mathbf{T}\mathbf{X}$  is the identity matrix  $\mathbf{I}$ .




# Whitening transformations

- We have thus derived a transform  $\mathbf{T} = \Lambda^{-\frac{1}{2}} \Phi^\top$  such that the (uncentered) auto-covariance of the transformed data  $\tilde{\mathbf{X}} = \mathbf{T}\mathbf{X}$  is the identity matrix  $\mathbf{I}$ .
- $\mathbf{T}$  transforms the cost from  $f_{\text{MSE}}(\mathbf{w}; \mathbf{X})$



# Whitening transformations

- We have thus derived a transform  $\mathbf{T} = \Lambda^{-\frac{1}{2}} \Phi^\top$  such that the (uncentered) auto-covariance of the transformed data  $\tilde{\mathbf{X}} = \mathbf{T}\mathbf{X}$  is the identity matrix  $\mathbf{I}$ .
- $\mathbf{T}$  transforms the cost from  $f_{\text{MSE}}(\mathbf{w}; \mathbf{X})$  to  $f_{\text{MSE}}(\mathbf{w}; \tilde{\mathbf{X}})$ :



# Whitening transformations

- Whitening transformations are a technique from “classical” ML rather than DL.
  - Time cost is  $O(m^3)$ , which for high-dimensional feature spaces is too large.
- However, whitening has inspired modern DL techniques such as **batch normalization** (Szegedy & Ioffe, 2015) (more to come later).

# Whitening transformations

- Demos (gradient\_descent, zscore).